

A Comparison of Neural Architectures for Movie Poster Genre Classification *

Samson LAM Boris LEUNG Johnny YIK Thomas TANG Grace LUK Michael LO

May 2024

Abstract

A movie poster is the first visual contact a viewer has with a film, and it signals genre before any other information reaches the audience. We measure how far that signal survives as a machine-readable one, and at what cost. Three regimes are compared, a convolutional network, four ImageNet backbones transferred to posters, and a vision-language model, against two backbones built for cheap inference, EfficientNetV2 and Fast-ViT. We found that the classifier draws on the whole poster rather than its center, tolerates rotation and relies mainly on color for edge detection. Poster design changes measurably across decades, and supplying the decade as an indicator raises the accuracy by about 4%. ResNets deliver the strongest results, reaching 54.5 % accuracy on the 19-genre task and 69.7 % on the simplified 3-genre task.

1 Introduction

A movie poster carries a deliberate visual argument. It is composed to tell a viewer, in the moment before any synopsis is read, what kind of film this is. Genre is the part of that argument the designer most wants to land, and it is encoded in choices a human reads without effort, the palette, the typography, the arrangement of faces, the presence or absence of a horizon. Whether those choices survive as a machine-readable signal is an empirical question, and it is the question this study addresses.

We treat genre recognition as image classification on the poster alone. The dataset is the Letterboxd collection published on Kaggle (Simon, 2024), from which 548,835 posters survive data cleaning, each labeled with one to four of 19 genres. This study compares different machine learning architectures, beginning with a convolutional neural network. VGG16 (Simonyan and Zisserman, 2015), AlexNet (Krizhevsky et al., 2012) and two depths of ResNet (He et al., 2016) supply the transfer learning. YOLOv8 (Jocher et al., 2023) contributes a detection backbone adapted to classification. CLIP (Radford et al., 2021) is evaluated zero shot, without any training on posters, and serves as the reference a fine tuned model has to beat. EfficientNetV2 (Tan and Le, 2021) and Fast-ViT (Vasu et al., 2023) are backbones designed for cheap inference, and they test whether the accuracy available at this task requires the parameter counts the older backbones carry.

The remainder of the report is organized as follows. Section 2 describes the dataset and the data cleaning pipeline. Section 3 reports a convolutional neural network trained. Section 4 covers the four transfer learning models. Section 5 covers the zero-shot and the fine-tuned efficient backbones. Section 6 places the comparison tables and concludes.

*Code and replication data are available at *GitHub*.

2 Data

The Letterboxd collection consists of 557,539 poster images with genre, country, language and release date tables. The image size of (345, 230, 3) dominates with 99.1 % of the data, as Table 19 in the appendix shows. We drop posters that do not conform to the dominant shape of (345, 230, 3), leaving 552,531. Next, movies carrying an implausible number of genres are treated as outliers. $[-0.904, 4.006]$ is the range of genre count within 3 standard deviations of the mean, so we set a ceiling of four genres per movie. This step discards a further 3,827 movies, leaving 548,835 movies.




Figure 1: Distribution of genres (left) and the number of genres per movie (right).

Figure 1 displays the shares of the 19 genres on the left, with Drama, Documentary and Comedy leading, and the number of genres carried on the right, where the mass sits on one or two. Figure 2 gives the production region and the spoken language. The dataset is dominated by European and North American productions and by English-language releases. This imbalance is used in Section 4.1, where we evaluate the trained classifier separately by region and language to test its accuracy dependence.




Figure 2: Distribution of regions (left) and languages (right).

The first release dates span from 1874 to 2031, with some posters created well in advance of the actual release date. Figure 3 illustrates their distribution by decade. On the grounds that the earlier material is sparse and stylistically distinct from the rest, we use only posters released after 1970, leaving a total of 424,102 posters after all data cleaning.




Figure 3: Movies by decade of first release.

3 Convolutional Neural Network

This section constructs a Convolutional Neural Network (CNN) from scratch, which fixes the floor every later model is measured against. Given the architecture limits, we remove two genres before training. Horror and Thriller rely on the same visual vocabulary, suspense, threat and psychological pressure, and are difficult to separate even by eye, so Thriller is folded into Horror. Western is also dropped due to a lack of images.

3.1 Model

We propose a 12-layer CNN model, with the initial 9 layers having 32 filters and the last 3 layers having 64 filters. All layers use a (3, 3) kernel, ReLU activation, and padding. Instance normalization is applied after each layer to improve learning, and max-pooling layers downsample the feature maps. After flattening, batch normalization is applied, and the data passes into a dense layer with 32 units and ReLU activation. Dropout regularization with a rate of 0.25 is then employed to prevent overfitting. In the next layer, batch normalization is applied again, and the data passes through another dense layer with 64 units and ReLU activation. Dropout regularization with a rate of 0.5 is applied at this stage. Finally, a dense layer with the number of genres G to classify and softmax activation is added, producing predicted probabilities for each genre. The model is compiled using the Adam optimizer, categorical cross entropy loss, and accuracy as the evaluation metric. Figure 14 in the appendix draws the stack.

The training set contains 9,000 posters: 7,200 for fitting and 1,800 for testing, balanced across

genres. An initial attempt with $G = 5$ genres produced accuracy, F1 and ROC curves no better than a random guess, which motivated the reduction to 3 genres reported below.

Poster design is not stationary. Figure 4 places one poster from each of four decades side by side, and the progression is visible without measurement. The 1910s poster is monochrome and relies on tonal contrast. The 1950s poster carries the vivid, broadly realistic palette Technicolor made available. By the 1990s, the colors are saturated and chosen for impact. The 2010s poster returns to restraint, using gradients and narrow tonal ranges to suggest depth. Therefore, we group movie poster release dates by decades, supplying this as an indicator.




Figure 4: Poster design across four decades, from the 1910s (Left) to the 2010s (right).

3.2 CNN Performance

We narrow our classification to 3 genres, namely [Action, Drama and Romance], then [Comedy, Drama and Romance], then [Drama, Horror and Music]. Table 1 reports the classification results under releases after 1970 and after 2000. Across repeated runs, the indicators raise accuracy by about 4%, which we take as confirmation that release period carries information the pixels do not supply on their own. Nevertheless, the performance is still disappointing, as it barely performs adequately on 3 genres. As a result, the remainder of the report turns to pre-trained models, where the representation is learned elsewhere, and only the classifier is fitted to posters.

Setting	AUROC	F1	Precision	Accuracy
Action, Drama, Romance, 1970 onward	60.2	42.6	42.9	44.5
Action, Drama, Romance, 2000 onward	66.2	46.6	45.4	44.9
Comedy, Drama, Romance, 1970 onward	62.0	45.6	46.3	47.3
Comedy, Drama, Romance, 2000 onward	62.1	45.6	45.6	46.5
Drama, Horror, Music, 1970 onward	62.5	45.2	45.3	46.6
Drama, Horror, Music, 2000 onward	53.0	41.7	42.4	41.4
Average	60.9	44.6	44.7	45.2

Table 1: Evaluation metrics for CNN, in percent.

4 Transfer Learning

Transfer learning reuses a representation learned on a large dataset and fits only a new classifier to the current task. The pre-trained body has already encoded edges, textures and shapes, so training is shorter, and the risk of overfitting a small sample is lower.

4.1 VGG16

4.1.1 Model

VGG16 is a 16-layer convolutional network built from a single repeated element, the 3×3 convolution (Simonyan and Zisserman, 2015). Its design argument is that a stack of small kernels reaches the receptive field of a larger one while using fewer parameters and inserting more rectifiers along the way, so depth buys expressiveness more cheaply than width. Convolutions are grouped into five blocks, each closed by a 2×2 max pool, which halves the spatial extent while doubling the channel count. Trained on ImageNet, the resulting body is a general-purpose feature extractor, and we use it here with frozen weights so only the classifier on top is fitted to posters.

We first continue with the 3 genre classification. After flattening, it feeds into a dense layer with 128 units and sigmoid activation, then dropout at 0.25, then a softmax layer with 3 outputs. Sigmoid is used in the single hidden layer because the network is shallow enough that saturation is not a concern. Table 2 gives the parameter counts.

Layer	Output shape	Parameters
Input	(345, 230, 3)	0
VGG16	(10, 7, 512)	14,714,688
Flatten	(35,840)	0
Dense, sigmoid	(128)	4,587,648
Dropout at 0.25	(128)	0
Output, softmax	(3)	387

Table 2: Layers of the VGG16 classifier with output shapes and parameter counts.

4.1.2 VGG16 Performance

Fitting minimizes categorical cross-entropy with Adam at a learning rate of 0.001, in mini-batches of 32, for 50 epochs. Figure 5 plots validation accuracy and loss by epoch on the left and the ROC curves on the right, and Table 3 gives the metrics on the held-out posters. Accuracy and F1 both sit near 61%, and the area under the curve reaches 79.0%, a clear advance over the CNN in Section 3 at a fraction of the training cost.

Model	AUROC	F1	Precision	Accuracy
VGG16	79.0	61.4	62.0	61.4

Table 3: VGG16 on the held-out 3-genre sample, in percent.




Figure 5: Fit and discrimination of the VGG16 classifier. Validation accuracy and loss by epoch on the left, ROC curves on the held-out sample on the right.

4.1.3 Model Inspection

While the VGG16 gives a moderately accurate result, we shall explain the classification procedures involved. We degrade the test posters in three controlled ways and re-run the trained classifier on each, holding everything else fixed and reading the F1 score. First, we inspect the importance of colors, corners and orientation. Figure 6 shows one poster under each transformation beside the original.



Figure 6: A test poster in its original form and under the three transformations, monochrome conversion, blanking of 30 border pixels and rotation by 180 degrees.

Table 4 gives the confusion matrices, each cell written as the original count followed by the count after the transformation.

Transformation	True genre	Predicted		
Monochrome, F1 = 59.4 %	Action	482 → 409	134 → 165	184 → 226
	Romance	119 → 80	510 → 507	150 → 192
	Thriller	169 → 133	170 → 177	482 → 511
Border removed, F1 = 45.7 %	Action	482 → 202	134 → 60	184 → 538
	Romance	119 → 63	510 → 244	150 → 472
	Thriller	169 → 45	170 → 40	482 → 736
Rotated, F1 = 54.9 %	Action	482 → 526	134 → 94	184 → 180
	Romance	119 → 247	510 → 345	150 → 187
	Thriller	169 → 252	170 → 115	482 → 454

Table 4: Confusion matrices under three transformations of the test posters, written as the original count followed by the transformed count. The original F1 is 61.4 %.

Removing color barely registers. We test how much the absence of color hinders our classification, since a more vibrant color scheme in movie posters could affect our judgment of their genre. Hence, we transformed the posters into black and white. However, averaging the three channels leaves the confusion matrix almost intact and F1 falls only to 59.4 %, from which we infer that the network uses color chiefly to locate edges, and that the grayscale gradient preserves those edges.

Removing the border does considerable damage. We test how much information the corners of the posters contain for the classification, since it seems logical that the center of the posters contains the most crucial information about the movies. Hence, we trimmed out 30 pixels from each corner, and surprisingly, it drives F1 down to 45.7 % and pushes most posters into the Thriller column. The classifier is therefore not concentrating on the center of the poster in the way a human viewer does. Its filters act across the whole image and any part removed is a loss.

Turning the poster upside down costs little. F1 holds at 54.9 % and the loss is concentrated in Romance. We interpret the results as Romance posters carry human contact more often than the other two genres, and the reading of a contact between two figures is sensitive to orientation in a way that a skyline or a weapon is not, which is a plausible account of why that genre alone suffers.

Next, we examine whether the classifier has cultural biases, since the dataset is dominated by Western and English-language releases. Tables 5 and 6 evaluate it separately by production region and by spoken language. Accuracy is lowest for African and South American releases, which also have the fewest examples. Across languages, the spread is narrower, and Hindi is the only clear shortfall, while Korean releases score above the English baseline. Oversampling underrepresented regions and languages can help, but we do not attempt it here.

	Africa	Asia	Europe	North America	South America	Oceania
F1	54.7	61.3	60.3	64.8	57.6	59.2

Table 5: VGG16 F1 by production region, in percent.

	English	Spanish	Cantonese	Hindi	Korean	Thai
F1	61.5	62.1	61.4	57.0	65.8	61.6

Table 6: VGG16 F1 by spoken language, in percent.

To gain insights into the operation, we analyze the feature maps. Reading the activations directly confirms the usual picture of a convolutional hierarchy. As the model progresses through its convolutional layers, it extracts more complex features from the images. Layer 8 responds to edges and layer 16 responds to shapes and finer structure, and on a poster for Ip Man it isolates the face. Figure 7 gives the extract.




Figure 7: Feature maps from layers 8 and 16 on a poster for Ip Man.

Overall, VGG16 is therefore rotation tolerant, largely indifferent to color and cultural origin, but reliant on the whole image rather than its center. However, it does not appear to capture the spatial relation between objects, which is where the deeper and more recent architectures in the following sections earn their place.

4.2 AlexNet and Residual Networks

While sections 3 and 4.1 reduced the problem to $G = 3$ genres, this section returns to the full 19-genre problem.

4.2.1 The Models

AlexNet is the network that made deep convolutional models the default for image classification, winning the 2012 ImageNet competition at a top-one accuracy of 63.3 % and by a wide margin over the runner-up (Krizhevsky et al., 2012). Five convolutional layers feed 3 fully connected ones, with rectifier activations throughout and dropout in the classifier. The implementation used here widens the channels to 64, 192, 384, 256 and 256 and carries 61.1 million parameters, almost all of them

in the two dense layers of width 4,096. By current standards, it is shallow and is included as the weakest of the pre-trained bodies rather than as a competitive candidate.

Residual networks (ResNets) answer the problem that depth alone creates. Past a certain number of layers, the gradient reaching the early ones vanishes, and accuracy degrades even on the training set. The remedy is to have each block learn a residual against its own input rather than a full transformation,

$$A^{(\ell)} = \varphi\left(\mathcal{F}\left(A^{(\ell-1)}; \theta^{(\ell)}\right) + A^{(\ell-1)}\right), \quad (1)$$

where $\mathcal{F}(\cdot; \theta^{(\ell)})$ is the stack of convolutions inside block ℓ , so the identity is available at no cost and the gradient has a path that skips the block entirely (He et al., 2016).

Figure 8 draws the two block types the two depths use. ResNet-18 stacks the basic block, two 3×3 convolutions of equal width. ResNet-50 stacks the bottleneck block, a 1×1 convolution that narrows the channels, a 3×3 convolution at that reduced width and a 1×1 convolution that widens by a factor of four, which buys depth at a fraction of the arithmetic. Table 7 sets out the three backbones. Figures 15 and 16 in the appendix draw the AlexNet stack and the residual stages. The pair of depths is what makes the comparison informative. If depth is what matters the two should separate, and if the residual connection is what matters both should stand clear of AlexNet.




Figure 8: The two residual blocks of Equation (1).

4.2.2 AlexNet and ResNet Performance

Both models take the poster resized to 224×224 . We replace each model’s softmax output with a linear layer with sigmoid activation, since a movie may carry several, and we use binary cross-entropy as the loss. Training runs for 5 epochs. We compare Adam and stochastic gradient descent, with the latter at a learning rate of 10^{-4} , while Adam uses 10^{-6} for AlexNet, 5×10^{-5} for ResNet-18, and 10^{-5} for ResNet-50. We track multi-label AUROC during training and validation.

	AlexNet	ResNet-18	ResNet-50
Stem	11×11 conv, stride 4	7×7 conv, stride 2	7×7 conv, stride 2
Body	five convolutions	2, 2, 2, 2 basic	3, 4, 6, 3 bottleneck
Widths	64 to 384	64 to 512	256 to 2,048
Head	two dense, 4,096	global pool, linear	global pool, linear
Parameters	61.1 M	11.7 M	25.6 M

Table 7: The three transfer learning backbones, as implemented in torchvision.

Adam converges faster and to a lower loss than stochastic gradient descent for both architectures. For AlexNet, the training losses are close, but the validation metrics separate clearly. For ResNets, the gap is even wider, as shown in Figure 9. Every result reported below uses Adam.




Figure 9: Validation AUROC by epoch under Adam and stochastic gradient descent, for AlexNet (left) and ResNet (right).

Figure 10 compares the three architectures under Adam. ResNet outperforms AlexNet throughout, consistent with greater depth allowing more abstract features and the skip connection in Equation (1) keeping the gradient usable through that depth.




Figure 10: Validation AUROC by epoch for AlexNet and ResNets under Adam.

Table 8 reports the 19-genre task and Table 9 the 3-genre task. ResNet clearly outperforms AlexNet, but the depth within the residual family does not improve significantly. ResNet-50 is far

deeper than ResNet-18, yet their training and validation curves nearly coincide, and ResNet-50 leads by a slim margin on every genre. The gap between AlexNet and ResNet is therefore attributable to the residual connection rather than to depth alone. Overall, ResNet-18 with shallower architecture performs slightly better in terms of evaluation metrics and training efficiency. Table 20 in the appendix gives AUROC by genre.

Model	AUROC	F1	Precision	Accuracy
AlexNet	54.5	13.4	22.7	42.9
ResNet-18	61.6	28.1	34.0	53.6
ResNet-50	62.6	29.9	33.6	54.5

Table 8: AlexNet and ResNet on 19-genre single-label classification, in percent.

Model	AUROC	F1	Precision	Accuracy
AlexNet	71.6	61.5	62.8	61.8
ResNet-18	77.4	69.8	69.7	69.7
ResNet-50	77.5	69.7	70.2	69.7

Table 9: AlexNet and ResNet on 3-genre single-label classification, in percent.

4.3 YOLOv8

4.3.1 The Model

YOLO, short for You Only Look Once, is a SOTA for real-time image/ video object detection rather than classification. It divides the image into a grid and predicts bounding boxes with class probabilities for each cell in a single forward pass, which makes it fast enough for video. The eighth release ships a classification branch, YOLOv8n-cls, in which the detection head is replaced by a classifier over the whole image (Jocher et al., 2023). Its backbone alternates five stride-two convolutions with four C2f blocks, each of which splits its input, passes one part through a chain of bottleneck units and concatenates the results, and the head is a global pool followed by a single linear layer. The nano variant used here is the smallest of five sizes, at 2.7 million parameters against the 61.1 million of AlexNet. It is included to test whether a backbone shaped by a detection objective, which must localize objects rather than summarize a scene, reads a poster differently from a classification backbone.

4.3.2 YOLOv8 Performance

The input convention is a directory tree, with posters filed under a folder named for their label and split into training and validation folders. We implement this using the pre-trained YOLOv8n-cls model in two ways. Under the multi-label approach, a poster is copied into one folder per genre

it carries. Under the multi-class approach, it appears once, under a single genre. Both train for 10 epochs with Adam and the default cross-entropy loss.

The multi-label approach barely learns. Over 10 epochs, the training loss falls by 0.2, and accuracy rises by 0.04, settling near 30%, and the normalized confusion matrix in Figure 11 shows that it separates nothing except Drama. Filing one poster under several conflicting labels appears to be the cause. The multi-class approach trains more cleanly, as Figure 12 shows.




Figure 11: Normalized confusion matrices, multi-label (left) and multi-class (right).




Figure 12: Top-one accuracy by epoch, multi-label (left) and multi-class (right).

Nevertheless, Table 10 shows the multi-label model ahead on both tasks despite its weaker training curve. A plausible reading is that the multi-class model learns one class deeply while the multi-label model acquires broader but shallower knowledge of the label space.

Task	Model	AUROC	F1	Precision	Accuracy
19 genres	YOLOv8, multi-label	82.8	29.7	29.7	29.7
	YOLOv8, multi-class	82.0	29.0	29.0	29.0
3 genres	YOLOv8, multi-label	60.4	39.6	39.6	39.6
	YOLOv8, multi-class	53.8	39.6	39.6	39.6

Table 10: The two YOLOv8 models on single-label classification, in percent.

5 Zero-Shot Prediction and Efficient Backbones

5.1 CLIP

CLIP is a multi-modal model trained to align images and text in a shared embedding space (Radford et al., 2021). Write f_I for the image encoder, f_T for the text encoder and

$$s_{ij} = \frac{f_I(X_i)^\top f_T(T_j)}{\|f_I(X_i)\| \|f_T(T_j)\|} \quad (2)$$

for the cosine similarity between the i th image and the j th caption of a batch of N matched pairs. Pre-training is contrastive. It maximizes the similarity of the N matched pairs while minimizing it across the $N^2 - N$ mismatched ones, by minimizing

$$\mathcal{L}_{\text{CLIP}} = -\frac{1}{2N} \sum_{i=1}^N \left[\log \frac{\exp(s_{ii}/\tau)}{\sum_{j=1}^N \exp(s_{ij}/\tau)} + \log \frac{\exp(s_{ii}/\tau)}{\sum_{j=1}^N \exp(s_{ji}/\tau)} \right], \quad (3)$$

with $\tau > 0$ a learned temperature. Because the two encoders share one embedding space, a classification task can be posed without any further training by writing each label as a sentence T_g and returning

$$\hat{y}(X) = \arg \max_{g \in \{1, \dots, G\}} \frac{f_I(X)^\top f_T(T_g)}{\|f_I(X)\| \|f_T(T_g)\|}, \quad (4)$$

as Figure 13 sets out. Figure 17 in the appendix places that step inside the contrastive pre-training of Equation (3).

That property makes CLIP the natural reference point for this study. It has never seen the poster dataset, so whatever it achieves is the accuracy available with no task-specific training, and a fine-tuned model has to beat it to justify its cost.




Figure 13: Zero-shot classification with CLIP. Each genre is written as one sentence, both sides are embedded, and Equation (4) returns the genre whose sentence lies closest to the poster.

The label set is written as the prompt *A movie poster of { } genre*, one sentence per genre, which is tokenized and passed through the text encoder. The poster is passed through the image encoder. The model returns a probability across the prompts, so an output of $[0.3, 0.3, 0.4]$ over Comedy, Documentary and Drama reads as a 0.4 probability of Drama. Four image encoders are evaluated, namely ResNet-50, a widened ResNet-50 at roughly four times the arithmetic, and the

base vision transformer at patch sizes 32 and 16, which stacks 12 transformer blocks of width 768 with 12 attention heads over the flattened patches (Dosovitskiy et al., 2021), drawn in Figure 18 of the appendix. The text encoder is held fixed across the four.

5.2 EfficientNetV2

EfficientNet scales a network along depth, width and input resolution together rather than one at a time (Tan and Le, 2019). Depth adds layers and so enlarges the receptive field, width adds feature maps, and resolution retains more of the image. The three are tied to a single compound coefficient ϕ ,

$$d = \alpha^\phi, \quad w = \beta^\phi, \quad r = \gamma^\phi, \quad \text{subject to} \quad \alpha\beta^2\gamma^2 \approx 2, \quad \alpha, \beta, \gamma \geq 1, \quad (5)$$

where d , w and r are the depth, width and resolution multipliers. The constraint is what makes the scaling economical, since the arithmetic grows linearly in depth but quadratically in width and in resolution, so a unit increase in ϕ raises the floating point operations by a factor of about two whichever way the budget is spent.

The second version adds progressive learning (Tan and Le, 2021), which begins training at a small image size under weak regularization and raises both as training proceeds, so that the expensive large-image passes occur only once the model is ready for them. Its body mixes two block types. The fused inverted bottleneck, which merges the expansion and the depthwise convolution into one ordinary convolution, is used in the early stages where the feature maps are large, and the standard inverted bottleneck is used later. Table 11 gives the stage structure.

5.3 Fast-ViT

Vision transformers are slow and memory inefficient. Splitting an image of resolution r into patches of side p yields a sequence of $n = (r/p)^2$ tokens, and self-attention costs $O(n^2)$ in that length, so halving the patch size multiplies the attention cost by a factor of 16. Fast-ViT is a hybrid design that keeps convolutions wherever a transformer would be expensive (Vasu et al., 2023). It is staged like a residual network, halving resolution and doubling channels from one stage to the next, it replaces patch embedding with a convolutional stem, and it uses convolutions in place of the linear feed-forward layers. That substitution rests on the MetaFormer observation that the transformer block keeps working when self-attention is replaced by another token mixer, a spatial multi-layer perceptron or even plain pooling (Yu et al., 2022), which Figure 19 in the appendix draws. Its own token mixer, RepMixer, uses a 3×3 depthwise convolution with batch normalization, reparameterized into a single convolution once training is complete. The variant used here is the tiny one with 12 blocks, the smallest of the family, in which every stage mixes tokens this way and no attention layer remains.

	EfficientNetV2-S	Fast-ViT T12
Stem	3×3 conv, 24 channels	3×3 convolutions, stride 4
Stage widths	24, 48, 64, 128, 160, 256	64, 128, 256, 512
Blocks per stage	2, 4, 4, 6, 9, 15	2, 2, 6, 2
Block type	fused and standard bottleneck	RepMixer
Head	1×1 conv to 1,280, linear	1×1 conv to 1,024, linear
Parameters	21.5 M	7.6 M

Table 11: The two fine tuned backbones of this section, as implemented in torchvision and in timm.

5.4 Fine Tuning with Time Embedding

Both models are fine-tuned in the same way. We load the pre-trained weights, freeze the body so only the head is estimated, and replace the head that outputs 1000 ImageNet classes with one that outputs the G genres. Write $z = \mathcal{A}(f_{\text{body}}(X)) \in \mathbb{R}^q$ for the globally averaged feature vector, with $q = 1,280$ for EfficientNetV2 and $q = 1,024$ for Fast-ViT. For Fast-ViT, the head is a single affine map to G outputs. For EfficientNetV2 it is two affine maps separated by dropout,

$$\eta = U_2 \text{Drop}_{0.2}(U_1 \text{Drop}_{0.2}(z) + c_1) + c_2, \quad \Pr(y = g \mid X) = \frac{\exp(\eta_g)}{\sum_{g'=1}^G \exp(\eta_{g'})}, \quad (6)$$

with U_1 of width 1,000. Training minimizes cross-entropy using Adam for either 10 or 22 epochs.

The time embedding variant adds the release year. The year enters as a one-hot vector $\iota(t) \in \{0, 1\}^K$ over the K distinct release years of the dataset, each coordinate of which indexes a learned four-dimensional embedding, so the year contributes a vector $u(t) \in \mathbb{R}^{4K}$. It is concatenated with the pooled features, so that the head of Equation (6) receives $(z^\top, u(t)^\top)^\top$ in place of z and U_1 widens to $q + 4K$ columns.

5.5 Performance

Both loaders were constructed over the entire frame, so the rows used for evaluation are the rows the models were fine-tuned on. Two label sets are reported, the full 19 genres and a reduced set of 3 genres [Comedy, Documentary and Drama]. Table 12 restricts the dataset by release period. Accuracy improves slightly as the window narrows, indicating a real but small time-varying element.

Image encoder	All periods	After 1970	After 2000
ResNet-50	25.5	25.7	26.4
ViT-B/32	31.7	33.6	34.4

Table 12: Zero-shot accuracy of CLIP over 19 genres by release period, in percent.

Tables 13 and 14 compare the four image encoders. The ordering is the same on both label sets and follows encoder strength, with the patch-16 vision transformer ahead throughout.

Image encoder	AUROC	F1	Precision	Accuracy
ResNet-50	75.9	16.7	19.3	25.5
ResNet-50 widened	78.2	18.4	20.4	25.5
ViT-B/32	79.3	19.0	20.4	31.7
ViT-B/16	81.3	21.0	22.0	32.9

Table 13: Zero-shot CLIP over 19 genres, in percent.

Image encoder	AUROC	F1	Precision	Accuracy
ResNet-50	74.4	52.8	55.8	53.2
ResNet-50 widened	71.6	45.4	51.2	51.0
ViT-B/32	74.3	51.3	56.3	55.1
ViT-B/16	79.8	60.0	60.5	60.9

Table 14: Zero-shot CLIP over Comedy, Documentary and Drama, in percent.

Tables 15 and 16 give the fine-tuned results against the best zero-shot figure. The time embedding hurts rather than helps EfficientNetV2 at both epoch counts, so we did not carry it over to Fast-ViT. On 19 genres, EfficientNetV2 leads Fast-ViT, plausibly because the 12-block Fast-ViT is the smallest variant of its family, and both lead the zero-shot reference, which is what fine-tuning is supposed to deliver. On 3 genres, zero-shot CLIP edges ahead of both at a fifth of the compute, which suggests that on an easy label set the textual description of the genre is worth more than task-specific fitting.

Model	Epochs	Without time embedding	With time embedding
EfficientNetV2	10		44.1
EfficientNetV2	22		30.4
Fast-ViT t12	10	36.9	
Fast-ViT t12	22	39.8	
CLIP, zero shot		32.9	

Table 15: Accuracy after fine tuning over 19 genres, against the best zero-shot encoder, in percent.

6 Conclusion

To conclude, Tables 17 and 18 display the performance of all model on the two tasks. We believe a movie poster carries enough genre signal to be worth modeling but not enough to classify a reliable label. Training a CNN from scratch is not worth the compute, since transfer learning achieves better results in a fraction of the time. We found that the drift in poster design is a real but small time-varying element in the data. On the model inspection, we discovered that a classifier is rotation tolerant, largely indifferent to color and cultural origin, but reliant on the whole image

Model	Epochs	Accuracy	Training and inference time
EfficientNetV2	22	59.2	135 minutes
Fast-ViT t12	22	60.8	100 minutes
CLIP, zero shot		60.9	20 minutes

Table 16: Accuracy after fine tuning against the best zero-shot encoder, in percent.

rather than its center, which is a different reading strategy from the one a human viewer brings to the same picture. Across the comparison, ResNets generally achieve the best results. YOLOv8 and ViT-B/16 also leads on area under the curve and demonstrate the capability of more complex and efficient architectures.

Model	AUROC	F1	Precision	Accuracy
CNN, average	60.9	44.6	44.7	45.2
VGG16	79.0	61.4	62.0	61.4
AlexNet	71.6	61.5	62.8	61.8
ResNet-18	77.4	69.8	69.7	69.7
ResNet-50	77.5	69.7	70.2	69.7
YOLOv8, multi-label	60.4	39.6	39.6	39.6
YOLOv8, multi-class	53.8	39.6	39.6	39.6
CLIP, ResNet-50	74.4	52.8	55.8	53.2
CLIP, ResNet-50 widened	71.6	45.4	51.2	51.0
CLIP, ViT-B/32	74.3	51.3	56.3	55.1
CLIP, ViT-B/16	79.8	60.0	60.5	60.9

Table 17: Evaluation metrics on 3-genre single-label classification, in percent.

Model	AUROC	F1	Precision	Accuracy
AlexNet	54.5	13.4	22.7	42.9
ResNet-18	61.6	28.1	34.0	53.6
ResNet-50	62.6	29.9	33.6	54.5
YOLOv8, multi-label	82.8	29.7	29.7	29.7
YOLOv8, multi-class	82.0	29.0	29.0	29.0
CLIP, ResNet-50	75.9	16.7	19.3	25.5
CLIP, ResNet-50 widened	78.2	18.4	20.4	25.5
CLIP, ViT-B/32	79.3	19.0	20.4	31.7
CLIP, ViT-B/16	81.3	21.0	22.0	32.9

Table 18: Evaluation metrics on 19-genre single-label classification, in percent.

References

- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *Proceedings of the International Conference on Learning Representations*, 2021.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- Glenn Jocher, Ayush Chaurasia, and Jing Qiu. Ultralytics yolov8. <https://docs.ultralytics.com/tasks/classify/>, 2023.
- Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems*, 2012.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, 2021.
- G. Simon. Letterboxd movies dataset. <https://www.kaggle.com/datasets/gsimonx37/letterboxd>, 2024.
- Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. In *Proceedings of the International Conference on Learning Representations*, 2015.
- Mingxing Tan and Quoc V. Le. Efficientnet: Rethinking model scaling for convolutional neural networks. In *Proceedings of the 36th International Conference on Machine Learning*, 2019.
- Mingxing Tan and Quoc V. Le. Efficientnetv2: Smaller models and faster training. In *Proceedings of the 38th International Conference on Machine Learning*, 2021.
- Pavan Kumar Anasosalu Vasu, James Gabriel, Jeff Zhu, Oncel Tuzel, and Anurag Ranjan. Fastvit: A fast hybrid vision transformer using structural reparameterization. In *Proceedings of the IEEE International Conference on Computer Vision*, 2023.
- Weihao Yu, Mi Luo, Pan Zhou, Chenyang Si, Yichen Zhou, Xinchao Wang, Jiashi Feng, and Shuicheng Yan. Metaformer is actually what you need for vision. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2022.

A Appendix

A.1 Poster Dimensions

Table 19 gives the distribution of image dimensions over the 557,539 posters of the raw download, from which Section 2 keeps only the dominant shape.

Image size	Count	Share (%)
(345, 230, 3)	552,531	99.1
(345, 230)	4,203	0.8
(344, 230, 3)	752	0.1
(344, 230)	32	0.0
Other sizes	21	0.0
Invalid format	1	0.0

Table 19: Poster dimensions with counts and shares of the corpus.

A.2 Model Architectures

Figure 14 sets out the network built for this report, where each box is one stage, labeled with the block it repeats, the number of repetitions and the width of its output in channels, and normalizations and activations are omitted. The remaining figures are reproduced from the papers that introduced those architectures. The VGG16 classifier is given as a layer table in Section 4.1 and the stage structure of EfficientNetV2 and Fast-ViT in Table 11, while the YOLOv8n-cls backbone is described in Section 4.3, so none of those is drawn again here.




Figure 14: The convolutional network built from scratch of Section 3.




Figure 17: CLIP, reproduced from Radford et al. (2021). Panel one is the contrastive pre-training of Equation (3) and panel three is the zero-shot rule of Equation (4), which is what this report uses.




Figure 18: The vision transformer, reproduced from Dosovitskiy et al. (2021). Two of the four CLIP image encoders are this architecture at patch sizes 32 and 16.




Figure 19: The MetaFormer block, reproduced from Yu et al. (2022). The token mixer on the left is a slot rather than a fixed choice, and the three instances to its right are attention, a spatial multi-layer perceptron and pooling. Fast-ViT places a depthwise convolution in that slot.

A.3 AUROC by Genre

Table 20 reports the per-genre AUROC behind the 19-genre results of Section 4.2.

Genre	AlexNet	ResNet-18	ResNet-50
Action	77.3	83.2	84.3
Adventure	72.1	77.1	78.6
Animation	87.9	94.5	95.2
Comedy	72.6	78.0	78.4
Crime	69.1	75.7	77.8
Documentary	75.9	83.4	83.7
Drama	66.6	74.6	75.0
Family	75.6	80.3	82.3
Fantasy	66.4	72.3	74.9
History	65.2	74.8	76.8
Horror	81.6	87.0	88.0
Music	67.6	78.2	80.0
Mystery	66.0	72.4	73.9
Romance	72.6	78.9	80.0
Science Fiction	72.5	78.6	81.0
TV Movie	70.0	78.2	80.1
Thriller	74.9	80.2	81.4
War	70.9	79.7	82.3
Western	83.8	92.1	93.3

Table 20: AUROC by genre for AlexNet and the residual networks, in percent.

A.4 Sample Posters by Decade

Figures 20 to 27 show 10 posters sampled from each decade. The pool is restricted to Hong Kong releases, and from 1980 onward to titles rated at or above three, which is how the original sample was drawn. They are therefore one narrow slice of the corpus rather than a cross section of it. They are shown because the drift in poster design that Section 3 measures is visible in them without measurement, from painted illustration and hand-set type in the earliest decades to photographic composition in the most recent.



Figure 20: Ten Hong Kong releases sampled from the 1950s.



Figure 21: Ten Hong Kong releases sampled from the 1960s.



Figure 22: Ten Hong Kong releases sampled from the 1970s.



Figure 23: Ten Hong Kong releases sampled from the 1980s.



Figure 24: Ten Hong Kong releases sampled from the 1990s.



Figure 25: Ten Hong Kong releases sampled from the 2000s.



Figure 26: Ten Hong Kong releases sampled from the 2010s.




Figure 27: Ten Hong Kong releases sampled from the 2020s.